


Sequential **Decisions** With Uncertainty: When Agents Learn

With slides from Pieter Abbeel and Dan Klein (Berkeley), and Percy Liang (Stanford)

HMMs for NLP: Tagging

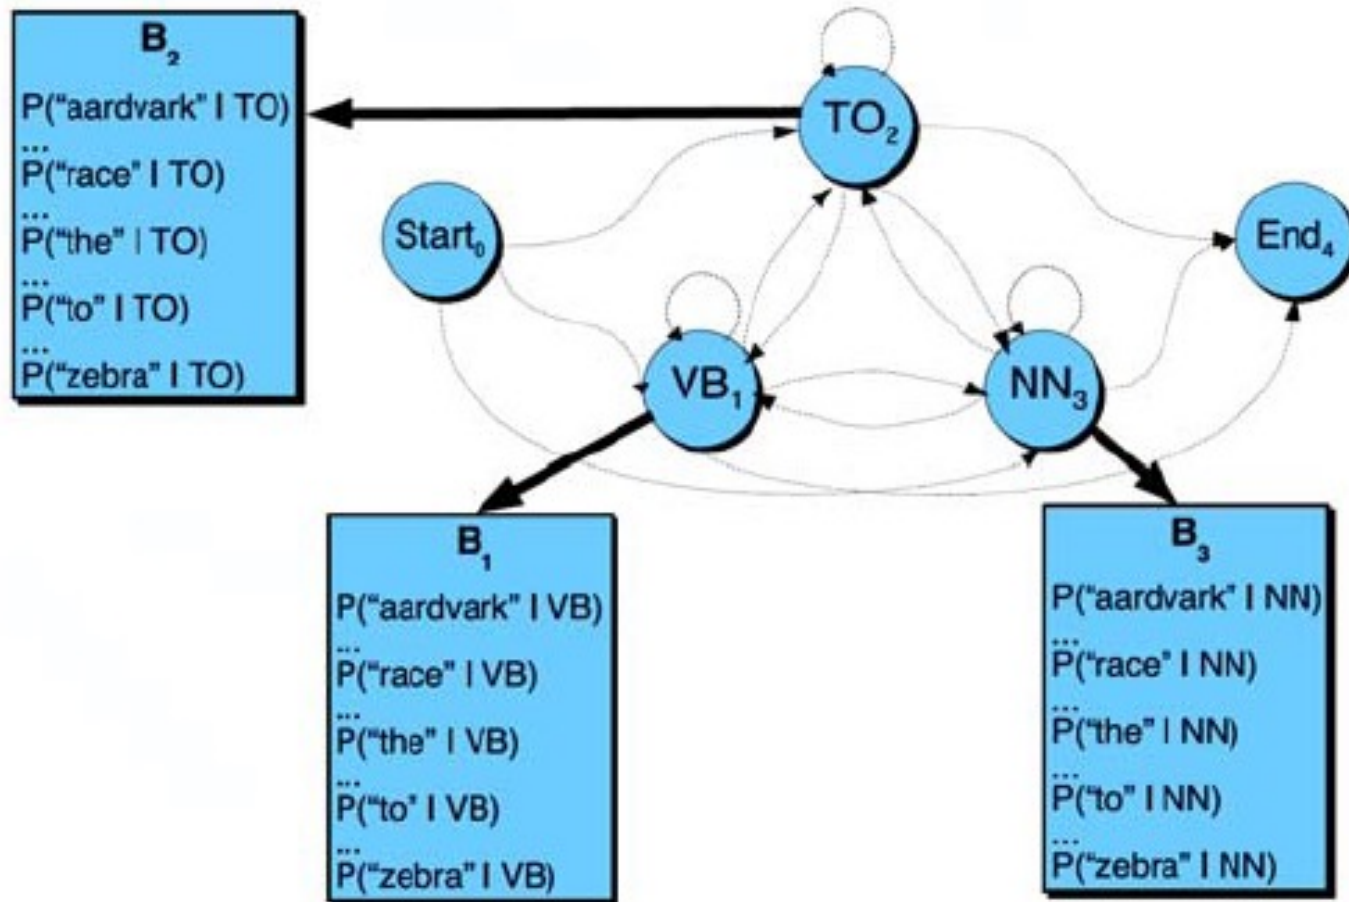
The Georgia branch had taken on loan commitments ...



DT NNP NN VBD VBN RP NN NNS


- HMM Model:
 - States $Y = \{DT, NNP, NN, \dots\}$ are the POS tags
 - Observations $X = V$ are words
 - Transition dist'n $q(y_i | y_{i-1})$ models the tag sequences
 - Emission dist'n $e(x_i | y_i)$ models words given their POS
- Q: How to we represent n-gram POS taggers?

Transitions and Emissions




Transitions

- Transitions $P(s|s')$ encode well-formed tag sequences
 - In a bigram tagger, states = tags



- In a trigram tagger, states = tag pairs



Inference (Decoding)

- Problem: find the most likely (Viterbi) sequence under the model

$$\arg \max_{y_1 \dots y_n} p(x_1 \dots x_n, y_1 \dots y_n)$$

- Given model parameters, we can score any sequence pair

NNP	VBZ	NN	NNS	CD	NN	.
Fed	raises	interest	rates	0.5	percent	.

$q(\text{NNP}|\blacklozenge) e(\text{Fed}|\text{NNP}) q(\text{VBZ}|\text{NNP}) e(\text{raises}|\text{VBZ}) q(\text{NN}|\text{VBZ}) \dots$

- In principle, we're done – list all possible tag sequences, score each one, pick the best one (the Viterbi state sequence)

NNP VBZ NN NNS CD NN  $\log P = -23$

NNP NNS NN NNS CD NN  $\log P = -29$


NNP VBZ VB NNS CD NN  $\log P = -27$

Entity Tagging with HMMs

Given a sequence of observations:


Yesterday Lawrence Saul spoke this example sentence.

and a trained HMM:



Find the most likely state sequence: (Viterbi)

$$\arg \max_{\vec{s}} P(\vec{s}, \vec{o})$$




Any words said to be generated by the designated “person name” state extract as a person name:

Person name: Lawrence Saul

HMM Example: “Nymble”

Task: Named Entity Extraction



Train on 450k words of news wire text.

Results:

Case	Language	F1 .
Mixed	English	93%
Upper	English	91%
Mixed	Spanish	90%

[Bikel, et al 1998],
[BBN “IdentiFinder”]

Transition
probabilities

$$P(s_t | s_{t-1}, o_{t-1})$$

Back-off to:

$$P(s_t | s_{t-1})$$

$$P(s_t)$$

Observation
probabilities

$$P(o_t | s_t, s_{t-1})$$

$$\text{or } P(o_t | s_t, o_{t-1})$$

Back-off to:

$$P(o_t | s_t)$$


$$P(o_t)$$

Remaining HMM/BN Issues

- Where do **P**, **E** probabilities come from?

- Training HMMs

- Dealing with unobserved values



- What's the most likely explanation (a.k.a. decoding)
 - Needed for speech recognition, language parsing, ...
 - Solution: Viterbi algorithm (most likely explanation)
 - What if we want multiple explanations? (beam search)

HMM Training: Degree of Supervision


- **Supervised**: Training data is tagged by humans
- **Unsupervised**: Training data isn't tagged
- **Partially supervised**: Training data isn't tagged, but you have a dictionary giving possible tags for each word
- We'll start with the supervised case and move to decreasing levels of supervision

Training an HMM (Supervised)

1. Define model topology (states, possible arcs)
2. Obtain labeled/tagged data
3. Estimate HMM parameters:
 1. Transition probabilities
 2. Emission probabilities
4. Validate on hold-out data
5. Done.

MLE Example: Squeezy Pig

- Consider a random pig:
 - Observations:
 - squeeze it 10 times, and 8 times it goes "oink" and 2 times it goes "wee"
 - Model topology:
 - assume it is a stateless pig (no hidden states)
 - Parameter estimation:
 - Maximum likelihood estimate (MLE) to model the pig as a random emitter with $P(\text{"oink"}) = 0.8$ and $P(\text{"wee"}) = 0.2$. $P(\text{"moo"})=?$
 - $P(x) = \text{count}(X) / \text{total number of trials}$



Learning: Maximum Likelihood

$$p(x_1 \dots x_n, y_1 \dots y_n) = q(STOP|y_n) \prod_{i=1}^n q(y_i|y_{i-1})e(x_i|y_i)$$

- Learning
 - Maximum likelihood methods for estimating transitions q and emissions e

$$q_{ML}(y_i|y_{i-1}) = \frac{c(y_{i-1}, y_i)}{c(y_{i-1})} \quad e_{ML}(x|y) = \frac{c(y, x)}{c(y)}$$

- Will these estimates be high quality?
 - Which is likely to be more sparse, q or e ?

Estimating Transitions: N-gram model

- Each word is predicted according to a conditional distribution based on a limited prior context
- Conditional Probability Table (CPT): $P(X|both)$
 - $P(of|both) = 0.066$
 - $P(to|both) = 0.041$
 - $P(in|both) = 0.038$
- From 1940s onward (or even 1910s – Markov 1913)
- a.k.a. Markov (chain) models

Estimating Emission Probabilities

$$P(s, w) = \prod_i P(s_i | s_{i-1}) P(w_i | s_i)$$

- Emissions are trickier:

- Words we've never seen before
- Will use new method: "smoothing"
- One option: break out the Good-Turning smoothing
- Issue: words aren't black boxes:

343,127.23 11-year Minteria reintroducibly

- Unknown words usually broken into word classes

D⁺, D⁺. D⁺ D⁺-x⁺ Xx⁺ x⁺"ly"

- Another option: decompose words into features and use a maxent model along with Bayes' rule

$$P(w | t) = P_{MAXENT}(t | w) P(w) / P(t)$$

Problem: Insufficient data


$$p(t \mid M_d) = 0$$

- Zero probability
 - May not wish to assign a probability of zero to an event that is never appeared in training data (unseen)
 - E.g., $p(\text{“moo”} \mid M_{\text{pig}})$
- General approach
 - An unseen event is possible, but no more likely than would be expected by chance.
 - How to assign $P(\text{unseen})$?

General Idea: Smoothing


- We often want to make estimates from sparse statistics:

$P(w \mid \text{denied the})$
3 allegations
2 reports
1 claims
1 request
7 total



- Smoothing flattens spiky distributions so they generalize better

$P(w \mid \text{denied the})$
2.5 allegations
1.5 reports
0.5 claims
0.5 request
2 other
7 total



- Very important all over NLP, but easy to do badly!
- We'll illustrate with bigrams today (h = previous word, could be anything).

Example: Laplace smoothing

- Idea: pretend we saw every word once more than we actually did [Laplace]
 - Corresponds to a uniform prior over vocabulary
 - Think of it as taking items with observed count $r > 1$ and treating them as having count $r^* < r$
 - Holds out $V/(N+V)$ for “fake” events
 - $N1+/N$ of which is distributed back to seen words
 - $N0/(N+V)$ actually passed on to unseen words (nearly all!)
 - Actually tells us not only how much to hold out, but where to put it
- Adding 1 **poorly** in practice (why?)
- Quick fix: add some small δ instead of 1 [Lidstone-Jefferys]
- Better better, holds out less mass

HMMs: Degree of Supervision

- **Supervised**: Training corpus is tagged by humans
- **Unsupervised**: Training corpus isn't tagged
- **Partly supervised**: Training corpus isn't tagged, but you have a dictionary giving possible tags for each word
- We'll start with the supervised case and move to decreasing levels of supervision

Motivation: Unsupervised Training

- Accuracy of tagging degrades outside of domain; Often make errors on important words (e.g., protein names)

$$\lambda = (A, B, \pi)$$

- We assumed that we know the underlying model
- Often these parameters are estimated on annotated training data, which has two drawbacks:
 - Annotation is difficult and/or expensive
 - Training data is different from the current data
- We want to maximize the parameters with respect to the current data, i.e., we're looking for a model λ' , such that

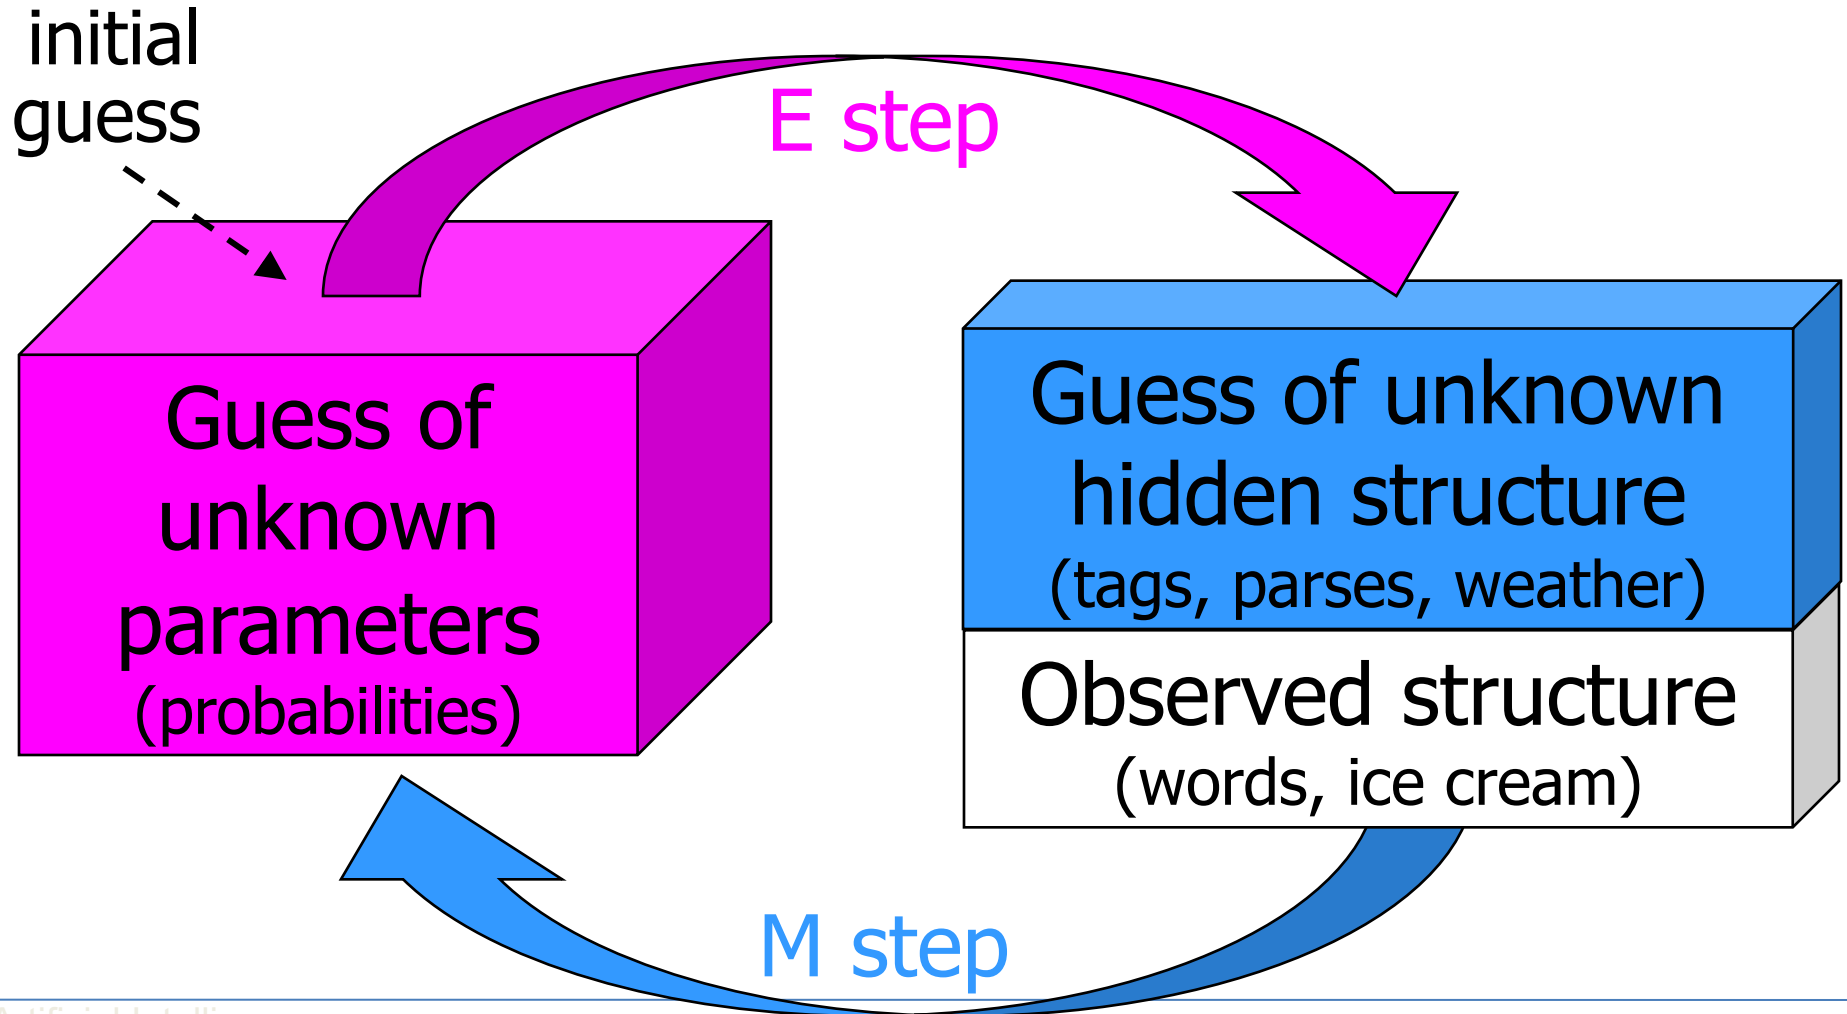
$$\lambda' = \underset{\lambda}{\operatorname{argmax}} P(O | \lambda)$$

General Idea

- Start by devising a noisy channel
 - Any model that predicts the corpus observations via some hidden structure (tags, parses, pigs...)
- Initially **guess** the parameters of the model!
 - Educated guess is best, but random can work
- **Expectation step:** Use current parameters (and observations) to reconstruct hidden structure
- **Maximization step:** Use that hidden structure (and observations) to reestimate parameters

Repeat until convergence!

General Idea



For Hidden Markov Models


initial
guess

E step

Guess of

	$p(\dots C)$	$p(\dots H)$	$p(\dots START)$
$p(1 \dots)$	0.7	0.1	
$p(2 \dots)$	0.2	0.2	
$p(3 \dots)$	0.1	0.7	
$p(C \dots)$	0.8	0.1	0.5
$p(H \dots)$	0.1	0.8	0.5
$p(STOP \dots)$	0.1	0.1	0

(probabilities)



M step

For Hidden Markov Models


initial
guess

E step

Guess of

	$p(\dots C)$	$p(\dots H)$	$p(\dots START)$
$p(1 \dots)$	0.677	0.058	
$p(2 \dots)$	0.219	0.425	
$p(3 \dots)$	0.105	0.517	
$p(C \dots)$	0.876	0.093	0.129
$p(H \dots)$	0.109	0.865	0.871
$p(STOP \dots)$	0.015	0.042	0

(probabilities)



M step

For Hidden Markov Models


initial
guess

E step

Guess of

	$p(\dots C)$	$p(\dots H)$	$p(\dots START)$
$p(1 \dots)$	0.697	0.04	
$p(2 \dots)$	0.171	0.464	
$p(3 \dots)$	0.132	0.496	
$p(C \dots)$	0.904	0.077	0.012
$p(H \dots)$	0.094	0.87	0.988
$p(STOP \dots)$	0.002	0.053	0

(probabilities)



M step

EM by Dynamic Programming: Two Versions

- The Viterbi approximation
 - Expectation: pick the *best* parse of each sentence
 - Maximization: retrain on this best-parsed corpus
 - *Advantage: Speed!*
- Real EM *why slower?*
 - Expectation: find *all* parses of each sentence
 - Maximization: retrain on *all* parses in proportion to their probability (as if we observed fractional count)
 - *Advantage: $p(\text{training corpus})$ guaranteed to increase*
 - Exponentially many parses, so don't extract them from chart – need some kind of clever counting

Baum-Welch (1)

- Want to estimate:
 - a_{ij} (transition probability from state i to state j)
- If we had tagged corpus:
 - $a_{ij} = \text{count}(\text{transition from state } i \text{ to state } j) / \text{count}(\text{transition from state } i \text{ to any state})$
- Don't have one. So:
 - **Assume some initial values** for a_{ij} and $b_j(w)$
 - Use a_{ij} and $b_j(w)$ to compute expected counts (E)
 - Then use E to *re-estimate* $a_{ij} \rightarrow a'_{ij}$
 - $a'_{ij} = E(\text{transition from state } i \text{ to state } j) / E(\text{transition from state } i \text{ to any state})$

BM: Intuition

- If training data has information about sequence of hidden states (as in word recognition example), then use maximum likelihood estimation of parameters:

$$a_{ij} = P(s_i | s_j) = \frac{\text{Number of transitions from state } S_j \text{ to state } S_i}{\text{Number of transitions out of state } S_j}$$

$$b_i(v_m) = P(v_m | s_i) = \frac{\text{Number of times observation } V_m \text{ occurs in state } S_i}{\text{Number of times in state } S_i}$$

BM: Approximation

General idea:

$$a_{ij} = P(s_i | s_j) = \frac{\text{Expected number of transitions from state } S_j \text{ to state } S_i}{\text{Expected number of transitions out of state } S_j}$$

$$b_i(v_m) = P(v_m | s_i) = \frac{\text{Expected number of times observation } V_m \text{ occurs in state } S_i}{\text{Expected number of times in state } S_i}$$

$$\pi_i = P(s_i) = \text{Expected frequency in state } S_i \text{ at time } k=1.$$

BM Algorithm: Iterations

- Recurrence:

1. Estimate $A_{k,l}$ and $E_k(b)$ from $a_{k,l}$ and $e_k(b)$ overall all training sequences (**E-step**)
2. Update $a_{k,l}$ and $e_k(b)$ using ML (**M-step**)
3. Repeat steps #1, #2 with new parameters $a_{k,l}$ and $e_k(b)$

- Initialization:

- Set A and E to pseudocounts (or priors)

- Termination: if $\Delta \log\text{-likelihood} < \text{threshold}$ or $N_{\text{times}} > \text{max_times}$

BM Algorithm: Iterations (2)

- **Recurrence:**

1. Calculate forward/backwards probs, $f_k(i)$ and $b_k(i)$, for each training sequence

2. **E-step:** Estimate the expected number of $k \rightarrow l$ transitions, $A_{k,l}$

$$A_{k,l} = \sum_i f_k(i) \cdot a_{k,l} \cdot e_l(x_{i+1}) \cdot b_l(i+1) / P(\vec{x} | \theta)$$

and the expected number of symbol b appearances in state k , $E_k(b)$

$$E_k(b) = \sum_{\{i | x_i = b\}} f_k(i) \cdot b_k(i) / P(\vec{x} | \theta)$$

3. **M-step:** Estimate new model parameters $a_{k,l}$ and $e_k(b)$ using ML across all training sequences
4. Estimate the new model's (log)likelihood to assess convergence



An example of Baum-Welch

(thanks to Sarah Wheelan, JHU)

- I observe the dog across the street. Sometimes he is inside, sometimes outside.
- I assume that since he can not open the door himself, then there is another factor, hidden from me, that determines his behavior.
- Since I am lazy, I will guess there are only two hidden states, S_1 and S_2 .





An example of Baum-Welch (cntd)

- Estimating initial probabilities:
 - Assume all sequences start with hidden state S_1 , calculate best probability
 - Assume all sequences start with hidden state S_2 , calculate best probability
 - Normalize to 1.
- Now, we have generated the updated transition, emission and initial probabilities. Repeat this method until those probabilities converge

An example of Baum-Welch (cntd)




- One set of observations:
 - I-I-I-I-I-O-O-I-I-I
- Guessing two hidden states. I need to invent a transition and emission matrix.
 - Note: since Baum-Welch is an EM algorithm the better my initial guesses are the better the job I will do in estimating the true parameters

Day $k+1$			
Day k		S1	S2
	S1	0.5	0.5
	S2	0.4	0.6

	IN	OUT
S1	0.2	0.8
S2	0.9	0.1

An example of Baum-Welch (cntd)



- Let's assume initial values of:
 - $P(S_1) = 0.3, P(S_2) = 0.7$
- Example guess: if initial I-I came from S_1 - S_2 then the probability is:
 $0.3 \times 0.2 \times 0.5 \times 0.9 = 0.027$

		Day $k+1$	
Day k		S1	S2
	S1	0.5	0.5
	S2	0.4	0.6

	IN	OUT
S1	0.2	0.8
S2	0.9	0.1



An example of Baum-Welch (cntd)

- Now, let's **estimate the transition matrix**. Sequence I-I-I-I-O-O-I-I-I has the following events:
 - II, II, II, II, IO, OO, OI, II, II
- So, our estimate for $S_1 \rightarrow S_2$ transition probability is:
 - $0.285 / 2.4474 = 0.116$
- Similarly, calculate the other three transition probs and normalize so they sum up to 1
- Update transition matrix

Seq	P(Seq) for $S_1 S_2$	Best P(Seq)
II	0.027	0.3403 $S_2 S_2$
II	0.027	0.3403 $S_2 S_2$
II	0.027	0.3403 $S_2 S_2$
II	0.027	0.3403 $S_2 S_2$
IO	0.003	0.2016 $S_2 S_1$
OO	0.012	0.0960 $S_1 S_1$
OI	0.108	0.1080 $S_1 S_2$
II	0.027	0.3403 $S_2 S_2$
II	0.027	0.3403 $S_2 S_2$
Total	0.285	2.4474

Summary – Baum-Welch Algorithm

- Start with initial HMM
- Calculate, using F-B, the likelihood to get our observations given that a certain hidden state was used at time i .
- Re-estimate the HMM parameters
- Continue until convergence
- Baum showed this to monotonically improve **data likelihood** (not necessarily accuracy!)
- Typically requires MANY observations, and (relatively) few parameters → simple topology

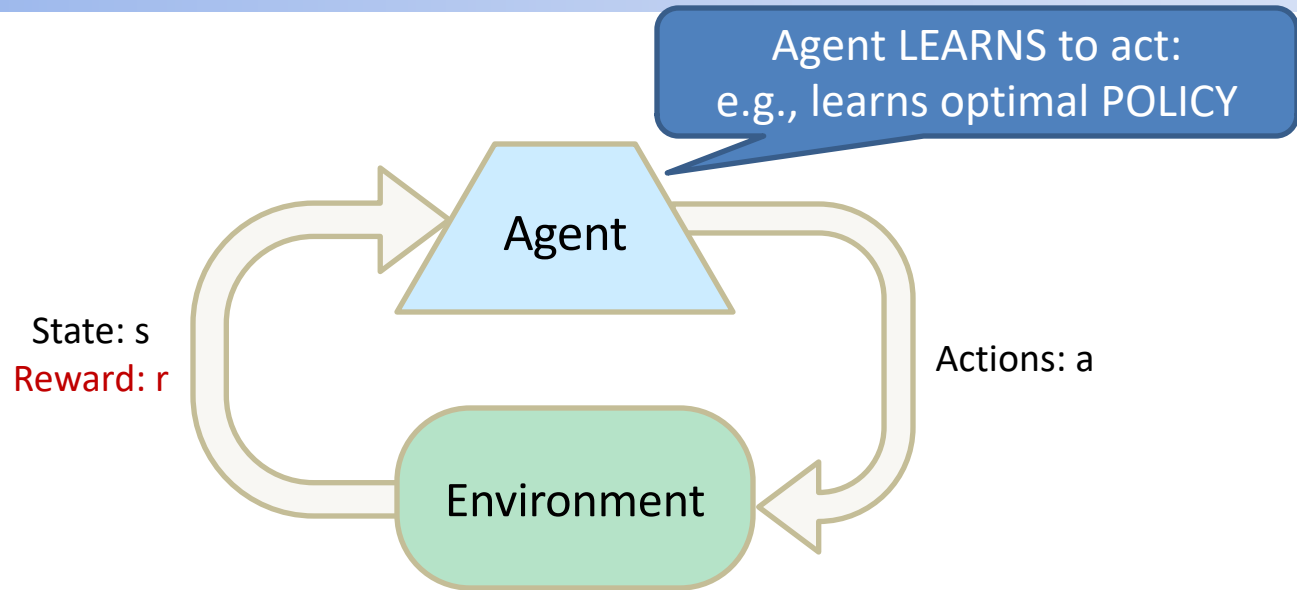
Summary: HMMs

- Widely used in AI: robotics, speech, NLP, ...
- Often have large number of hidden states (e.g., words in vocabulary \rightarrow 500K). Optimization needed
- Training exhaustively using MLE not feasible, need to generalize to unseen input (smoothing, classes)
- Now: **Markov Decision Process**.
 - How to make sequential decisions to optimize total expected future rewards

Big Picture

- AI as Planning:
 - ✓ Model of the world known (utilities, action outcomes)
 - ✓ Deterministic search: UCS, A*, MiniMax
 - ✓ Non-deterministic search → ExpectiMax
 - Inference under uncertainty: BNs, HMMs
- AI as Learning:
 - Model of world *partially* known (rewards? outcomes?)
 - Markov Decision Processes (Today)
 - Rewards, action outcomes unknown
 - Reinforcement Learning (After Spring Break)

Big Picture: Learning Agents



- Basic idea:
 - Receive feedback in the form of **rewards**
 - Agent's utility is defined by the reward function
 - Must (learn to) act so as to **maximize expected rewards**
 - All learning is based on observed samples of outcomes!


Rewards, Transitions known: MDP

Rewards, Transitions not known: ????

Plan (Next few weeks)


- Markov Decision Processes (MDPs)
 - MDP formalism
 - Solution: Value Iteration and Policy Iteration
- Reinforcement Learning (RL)
 - Relationship to MDPs
 - Several learning algorithms
 - RL applications to games, “real world”
- Project 3: MDPs, RL for Pacman and gridworld

Non-Deterministic **Actions**




Example: Grid World

- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- **Noisy movement:** actions do not always go as planned
 - 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
 - Small "living" reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize the sum of rewards




Grid World Actions

Deterministic Grid World




Stochastic Grid World




New Idea: Markov Decision Process (MDP)

- An MDP is defined by:
 - Set of states $s \in S$
 - Actions $a \in A$
 - Transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - Reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - Start state (s_0)
 - Terminal state (optional)
- MDPs are non-deterministic search problems
 - One way to solve them is with **expectimax** search
 - We can do better



What is “Markov” about MDPs?


- Remember: “Markov” means that **given the current state**, the **future** and the **past** are **independent**



Andrey Markov
(1856-1922)

- Like (H) MMs, where **the successor function only depends on current state** (not the full history)

What is “Markov” about MDPs (2)?



- Markov decision processes, “Markov” means action outcomes **depend only on the current state**

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$$


=

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$

- Like (H) MMs, where **the successor function only depends on current state (not the full history)**


New Idea: Markov Decision Process (MDP)

- An MDP is defined by:
 - Set of states $s \in S$
 - Actions $a \in A$
 - Transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - Reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - Start state (s_0)
 - Terminal state (optional)
- MDPs are non-deterministic search problems
 - One way to solve them is with **expectimax** search
 - We can do better




MDP Search Trees

- Each MDP state can be viewed as an **expectimax search node**




Expectimax Solution: Best Action from S



Definition: Policy


- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal **policy** $\pi^*: S \rightarrow A$
 - A policy π gives an **action** for each state
 - An **optimal policy** is one that **maximizes expected utility** if followed
 - An explicit policy defines a reflex agent
- Minimax/Expectimax **do NOT** compute entire **policies**
 - **They computed the action for a single state only (and then re-plan)**




Optimal policy when
 $R(s, a, s') = -0.03$ for all
non-terminals s

Optimal Policies


$R(s)$ = the “living reward”




$R(s) = -0.01$



$R(s) = -0.03$



$R(s) = -0.4$



$R(s) = -2.0$

Utilities of Sequences


- What preferences should an agent have over reward sequences?

$[1, 2, 2]$ or $[2, 3, 4]$

- More or less?


$[0, 0, 1]$ or $[1, 0, 0]$

- Now or later?



Infinite Utilities?!

- Problem: What if the game lasts forever? Do we get infinite rewards?
- Solutions:
 1. Finite horizon: (similar to depth-limited search)
 - Terminate episodes after a fixed T steps (e.g. assume finite lifetime)
 - *Problem*: Gives nonstationary policies (π depends on time left)
 2. Discounting: use $0 < \gamma < 1$
$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max}/(1 - \gamma)$$
 - Smaller $\gamma \rightarrow$ smaller “horizon” – shorter term focus
 3. Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like “overheated” for racing)



Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution: values of rewards decay exponentially



1

Worth
Now



γ

Worth Next
Step




γ^2

Worth In Two
Steps

Big Idea: Reward Discounting

- How to discount?
 - Each time we descend a level, we multiply by the discount once
- Why discount?
 - Sooner rewards probably do have higher utility than later rewards
 - Also helps our algorithms converge
- Example: discount of 0.5
 - $U([1,2,3]) = 1*1 + 0.5*2 + 0.25*3$
 - $U([1,2,3]) < U([3,2,1])$



Detour: Temporal/Delay Discounting

- What would you rather have?
 - A. \$100 today
 - B. \$150 a year from now
- What about:
 - A. \$100 in 12 months
 - B. \$110 in 13 months
- Humans temporally discount values of rewards
 - https://en.wikipedia.org/wiki/Temporal_discounting
- Delayed gratification:
https://www.youtube.com/watch?v=QX_oy9614HQ

Quiz: Discounting


- Given:

10				1
a	b	c	d	e

- Actions: East, West, and Exit (only available in exit states a, e)
 - Transitions: deterministic (no noise, for now)
- P 1: For $\gamma = 1$, what is the optimal policy?
 - P 2: For $\gamma = 0.1$, what is the optimal policy?
 - P 3: For which γ are West and East equally good when in state d?


<https://www.wolframalpha.com/input/?i=10x%5E3+%3D+x>

Example: Car Racing




Car Racing: State Diagram

- A robot car wants to travel far, quickly
- Three states: **Cool**, **Warm**, Overheated
- Two actions: *Slow*, *Fast*
- Going faster gets double reward




Racing Search Tree (depth=?)



Recap: Defining MDPs

- Markov decision processes:
 - Set of states S
 - Start state s_0
 - Set of actions A
 - Transitions $P(s' | s, a)$ (or $T(s, a, s')$)
 - Rewards $R(s, a, s')$ (and discount γ)
- MDP quantities so far:
 - **Policy** = Choice of action for each state
 - **Utility** = sum of (discounted) rewards



Next: Solving MDPs

